

SDN-Based Stateless Firewall Project

Purpose

This project allows learners to delve into Software-Defined Networking (SDN) principles by implementing a firewall using OpenFlow rules. They gain practical experience in customizing and enforcing security policies and monitoring traffic. Through this project, learners will understand scalability, performance considerations, and risk assessment in network security, enhancing their understanding of both SDN and network security concepts.

Objectives

Learners will be able to:

- Understand Software Defined Network (SDN) and its architecture better
- Setup Openflow environment, flow rules, and do packet processing using Openflow
- Setup a simple Open Virtual Switch (OVS) network using miniment or containernet
- Setup POX OpenFlow controllers
- Implement flow-based filtering rules based on flow control requirements.
- Implement flow-based firewall testing

Technology Requirements

- Hardware
 - Intel or AMD based computer with 8GB or more memory.
 - NOTE: 6GB is technically sufficient, but performance will be sluggish.
- Software
 - [VirtualBox 5.0](#) or Newer
 - Windows 10, Mac OS X, or Linux 64-bit as the base operating system.
 - Access to Apporto through a web browser

- Linux: Ubuntu 18.04 LTS (Bionic Beaver)

Project Description

In this project, you will emulate Denial of Service (DoS) attacks, which can target various components in the SDN infrastructure, in an SDN networking environment. You will need to set up an SDN-based firewall networking environment based on mininet, containernet, POX controller, and Over Virtual Switch (OVS). To mitigate DoS attacks, you will need to develop a mitigation strategy by using an SDN-based firewall to counter the implemented DoS attacks.

To complete the lab, you will submit a sequence of screenshots and corresponding illustrations to demonstrate how they fulfilled the firewall's packet filtering requirements. Your work will be evaluated based on the completeness of implementing firewall filtering rules to demonstrate how to successfully counter the implemented DoS attack.

Directions

You may work on your project locally or through Apporto. Follow the access guidelines provided and then review the Project Preparation steps to complete your work. Headings or titles labeled with "Local only" are specific to the local setup, otherwise you may complete your work through Apporto as well. Please make sure your final deliverable follows the requirements.

Project Files

The project files include the VM image file, the assigned project, and associated background labs. These can be found in your course on the Project Overview page.

Local Setup

Setup your local environment based on the directions provided below:

1. Download VirtualBox 5.0 or newer. **Note:** This step is not required if you have installed the VirtualBox in the Packet Filter Firewall (iptables) Project.
2. Download the Virtual Machine Image File mentioned in the Project Files section

3. Extract the .vdi file from the downloaded zip file. Note that the image sizes are large and sometimes you may not be able to decompress to get your image file. In such case follow the instructions below:
 - a. Go to your terminal and do: `unzip file_name -d /path/to/directory`
4. After extracting the image files, follow the instructions in the background lab **CS-SYS-00101 (VM in VirtualBox)** available in the **Associated Background Labs** zip file to set up the Virtual Machine.
5. Server credentials:
 - a. Username: ubuntu
 - b. Password: 123456

Accessing Apporto

For this project, you will work through Apporto's modular cyberlabs available through the course. Review all the project directions before starting so you know how to submit your work correctly:

1. In the course, click "Apporto Virtual Lab – App Store" (be patient while the lab loads)
2. In the App Store, select the "**CSE 548 Modular CyberLab**" lab. Please be patient as the lab initializes, which may take a few moments.
3. Once the lab environment is ready, you will be on a Linux machine and you need to click on the "**GNS3**" app from the Activities Menu.
4. In the pop-up window, choose "Open a project from disk". Navigate to the path Home → Labs → projects → CSE548 Project 2 and then select the "**CSE548 Project 2.gns3**" file.
5. You will see two (2) components listed: a Linux Server(Project2) and a NAT Network(NAT1) with an IP range of 10.0.2.0/24.
6. Right-click on the Linux server and select "**Start**", to start the server. You will know it's running when the red block next to the server changes to green.
7. Again, right-click on the server and select "**console**" to access the Linux desktop.
8. Server credentials:
 - a. Username: ubuntu
 - b. Password: 123456
9. You are all set! You can now start working on your project.

Note: Learners should refrain from using the virtual machine (VM) for personal purposes. This VM is provided solely for academic and course-related activities. Any personal activities may interfere with its intended purpose and impact the learning environment for yourself and your peers.

Project Preparation:

Assigned Project and Associated Background Labs

- Unzip the zip file for the labs using: `unzip file_name -d /path/to/directory`

Background Labs

Associated Background Labs

- CS-NET-00006 (OpenVirtual Switch);
- CS-NET-00007 (Mininet);
- CS-NET-00008 (POX Controller);
- CS-NET-00009 (Containernet)

Task 1.0 Preparation of setting up lab environment

Suggestions for Task 1:

1. Review and exercise the following labs CS-SYS-00001 (Linux tutorial), CS-NET-00001 (Network setup), and CS-NET-00002 (Gateway setup) before you do Task 1.1.
2. Review and exercise the following labs: Open vSwitch and Basic Setup (CS-NET-00006), mininet (CS-NET-00007), POX (CS-NET-00008), and containernet (CS-NET-00009) before you do Task 2.1.

Note that the services (Mininet, OVS, and POX) may have already been set up on your server. You need to verify if these servers are set up properly to conduct your required tasks.

Project Directions:

Before starting the lab, you need to check several software components and see if they have been set up properly. They are:

- Check if Python is installed

```
$ python --version
```

- Check if python3.x is installed

```
$ python3 --version
```

- check if mininet is installed (or check CS-NET-00007 for mininet installation and setup)

```
$ mn --version
```

```
$ sudo mn --test pingall
```

This will create temporary hosts and switch and pings to all hosts that it created. If mininet is installed properly, it will execute and clean all the created hosts later and exit successfully.

- Check installation of pox. Go to the directory where POX is installed, e.g., a common source code of POX is installed in the directory /home/ubuntu/pox. Here, we use \$POX DIR to represent the POX directory in your system.

```
$ cd /home/ubuntu/pox # depends on where the pox folder is created.  
$ ./pox.py -verbose forwarding.hub
```

This will start a pox session and print DEBUG messages as presented below. This indicates that POX is running fine. To exit press Ctrl-C.

```
POX 0.5.0 (eel) / Copyright 2011-2014 James McCauley, et al.  
INFO:forwarding.hub:Proactive hub running.  
DEBUG:core:POX 0.5.0 (eel) going up...  
DEBUG:core:Running on CPython (2.7.17/Apr 15 2020 17:20:14)  
DEBUG:core:Platform is Linux-5.3.0-40-generic-x86_64-with-Ubuntu-18.04-bionic  
INFO:core:POX 0.5.0 (eel) is up.  
DEBUG:openflow.of_01:Listening on 0.0.0.0:6633
```

- Check OVS installation

```
$ ovs-vsctl --version
```

It should display OVS version and details when it was installed, something like the below:

```
ovs-vsctl (Open vSwitch) 2.9.5  
DB Schema 7.15.1
```

Task 2.0 Test OpenFlow-based SDN Firewall

A firewall is used as a barrier to protect networked computers by blocking malicious network traffic generated by attackers. Internal traffic is not seen and cannot be filtered by a traditional firewall. An

SDN firewall works both as a network traffic flow filter and a policy checker. An example of an OpenFlow-based SDN setup is shown in Figure CS-CNS-00101.2. The first packet of a flow goes through the controller and is checked by the SDN controller based on its allowed or disabled network traffic policy, in which the subsequent packets of the flow directly match the flow policy defined in the controller. The firewall policy is centrally defined and enforced at the controller.

Security policies for a network are defined centrally at the controller. The firewall converts these policies into flow rules which are installed in the network by flow programming application. Firewall will have a centralized view of a network and it can inspect traffic to detect and drop unwanted network traffic. Unlike traditional firewalls that scan and filter every packet coming and going through it, an SDN-based firewall analyzes the first packet of each flow and installs rules for the rest of that flow in relevant switches. OpenFlow provides a mechanism to communicate to the controller from switches and vice-versa. A firewall application based on OpenFlow 1.0 and POX controller is used in this lab to demonstrate lab work in the SDN environment.

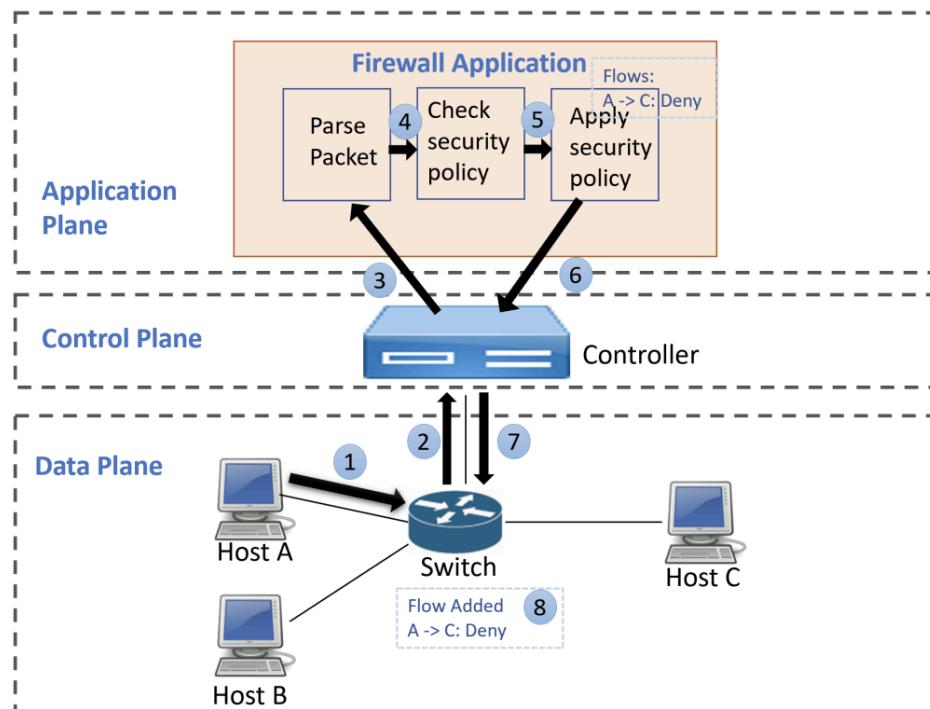


Figure CS-CNS-00101.2
Network topology for SDN.

OpenFlow

The control plane in SDN makes the forwarding decisions and applies policies to the entire network. This information sharing between the SDN controller and switches takes place using APIs provisioned by OpenFlow protocol. OpenFlow-enabled switches maintain the flow routing policies in flow tables. A typical rule presented in a flow table has 3 different field sets for packet handling:

- Match,
- Action, and
- Statistics

The match set allows packet filtering based on header fields. After a successful match, the packet undergoes respective actions. As per the OpenFlow standard, the first packet of a network flow is typically sent to the controller, which then inspects the packet headers and decides actions to be taken for the rest of the flow. This action can be either to drop or forward. The central visibility in SDN is of great advantage as the controller can extrapolate the future impact of traffic on any other node in the network.

SDN Controller

In an SDN, a controller acts as the heart of the network operating system. Hardware-based services of traditional networks like firewalls and load balancers run as software applications within a controller.

Task 2.1 Getting source code

1. Now, you can download lab resource files by using the following command. (Pre-requisite: Check if wget is installed using the command 'wget –version'. If wget is not installed, install it using 'sudo apt install wget')

```
$ wget
https://raw.githubusercontent.com/SaburH/CSE548/main/lab-cs-cns-00101
.zip
$ unzip lab-cs-cns-00101.zip
```

Source code files are located in the folder 'lab-cs-cns-00101'.

2. We need to use this firewall application with a POX controller. To do so, copy the L3Firewall.py file into the pox folder such as ./pox/pox/forwarding. Here we assume POX DIR is the directory where POX source code is present, e.g., If you have installed POX in /home/ubuntu/, then your \$POX DIR is /home/ubuntu/pox.

```
$ sudo cp lab-cs-cns-00101/l2firewall.config $POX_DIR/. # copy the
layer 2 config file.
$ sudo cp lab-cs-cns-00101/l3firewall.config $POX_DIR/. # copy the
layer 3 config file.
$ sudo cp lab-cs-cns-00101/L3Firewall.py
$POX_DIR/pox/forwarding/. # copy the forwarding file.
```

First, let's dive into the example L3Firewall application, please refer to Figure CS-CNS-00101.2 as the reference example for the communication flow between hosts, controller, and switch.

POX application starts off with launching the class defined with additional parameters, if any. In this case, the codes register the class “learning” with input configuration files in “l2config” and “l3config”. Users can provide these files using options –l2config and –l3config, if not provided, the program takes the default files provided in the application.

```
def launch (l2config="l2firewall.config",l3config="l3firewall.config"):
    parser = argparse.ArgumentParser()
    parser.add_argument('--l2config', action='store', dest='l2config',
                        help='Layer 2 config file', default='l2firewall.config')
    parser.add_argument('--l3config', action='store', dest='l3config',
                        help='Layer 3 config file', default='l3firewall.config')
    core.registerNew(Firewall,l2config,l3config)
```

The learning class consists of basic input parameters to start the application. Here, l2config and l3config file consists of user-specified firewall rules. That will get added in the fwconfig set which will contain individual rules.

Now, this initialization calls handle ConnectionUp as part of POX initialization architecture. POX connection is set here and basic layer 2 rules are installed in the firewall.

```
def _handle_ConnectionUp (self, event):
    self.connection = event.connection
    for (source, destination) in self.disabled_MAC_pair:
        message = of.ofp_flow_mod() # OpenFlow massage. Instructs a switch to
        install a flow
        match = of.ofp_match() # Create a match
        match.dl_src = source # Source address
        match.dl_dst = destination # Destination address
        message.priority = 65535 # Set priority (between 0 and 65535)
        message.match = match
        event.connection.send(message) # Send instructions to the switch
    log.debug("Firewall rules installed on %s", dpidToStr(event.dpid))
```

All incoming packets are handled by the handle PacketIn function. This will define what to do with incoming packets. From the incoming packet flow, match fields are retrieved and packets

are handled based on their protocol type, e.g., APR packets are handled differently than IP packets.

```
def _handle_PacketIn(self, event):
    packet = event.parsed
    match = of.ofp_match.from_packet(packet)

    if(match.dl_type == packet.ARP_TYPE and match.nw_proto == arp.REQUEST):
        self.replyToARP(packet, match, event)

    if(match.dl_type == packet.IP_TYPE):
        ip_packet = packet.payload
        print "ip_packet.protocol = ", ip_packet.protocol
        if ip_packet.protocol == ip_packet.TCP_PROTOCOL:
            log.debug("TCP it is !")
        self.replyToIP(packet, match, event, self.rules)
```

Once the packets are identified and separated according to protocols, they are sent to InstallFlow. This is the main part of the code where the OpenFlow rule is formed based on rules in the config file provided.

There are two types of timer values used for the OpenFlow flows. The idle timeout value determines how long a flow rule can be present in the flow table if there is no network traffic associated with that flow. In this case, for each flow, if there is no packets hitting a rule, that rule is removed from the device. Hard timeouts are the absolute values of time after which the rule is removed from the device. The switch maintains an entry time for each flow rule. The default timeout value for a flow is zero unless specified. In that case, an entry remains in the table permanently unless explicitly deleted.

```
def installFlow(self, event, offset, srcmac, dstmac, srcip, dstip, sport, dport, nwproto):
    msg = of.ofp_flow_mod()
    match = of.ofp_match()
    if(srcip != None):
        match.nw_src = IPAddr(srcip)
    if(dstip != None):
        match.nw_dst = IPAddr(dstip)
    match.nw_proto = int(nwproto)
    match.dl_src = srcmac
    match.dl_dst = dstmac
```

```

match.tp_src = sport
match.tp_dst = dport
match.dl_type = pkt.ethernet.IP_TYPE
msg.match = match
msg.hard_timeout = 0
msg.idle_timeout = 7200
msg.priority = priority + offset
event.connection.send(msg)

```

The input configuration file l2firewall.config file contains the following rules for each line in the configuration file:

[Priority, Source_MAC, Destination_MAC](#)

The configuration file is specified as:

```

id,mac_0,mac_1
1,00:00:00:00:00:01,00:00:00:00:00:02

```

The input configuration file l3firewall.config file contains the following bi-directional rule fields for each line in the configuration file:

[Priority, Source_MAC, Destination_MAC, Source_IP, Destination_IP,
Source_PORT, Destination_PORT, Network_PROTOCOL](#)

The configuration file is specified as:

```

priority,src_mac,dst_mac,src_ip,dst_ip,src_port,dst_port,nw_proto
1,any,any,10.0.0.1,10.0.0.2,1,1,icmp
2,any,any,10.0.0.2,10.0.0.1,1,1,icmp
3,any,any,10.0.0.1,10.0.0.2,any,any,tcp
4,any,any,10.0.0.2,10.0.0.1,any,any,tcp
5,any,any,10.0.0.1,10.0.0.2,any,any,udp
6,00:00:00:00:00:03,any,10.0.0.3,any,any,any,tcp

```

Note that POX takes the default IP ranges from “10” network if no specified IP address is given for each host, in which it assigns 10.0.0.1 to h1 . . . 10.0.0.9 to h9. This firewall configuration file consists of firewall rules for BLOCKING traffic. In l3firewall.config file:

[The first rule BLOCKS ICMP packets from source IP address 10.0.0.1 to](#)

destination IP address 10.0.0.2 and vice versa

The second rule BLOCKS ICMP packets from source IP address 10.0.0.2 to destination IP address 10.0.0.3 and vice versa

The third rule BLOCKS TCP packets from source IP address 10.0.0.1 to destination IP address 10.0.0.2 and vice versa

The fourth rule BLOCKS TCP packets from source IP address 10.0.0.2 to destination IP address 10.0.0.1 and vice versa

The fifth rule BLOCKS UDP packets from source IP address 10.0.0.1 to destination IP address 10.0.0.2 and vice versa

The sixth rule BLOCKS TCP packets from source MAC address 00:00:00:00:00:03 and IP address 10.0.0.3 and vice versa

3. Run POX controller, where here we assume \$POX DIR is the POX directory where pox source codes present:

```
$ cd $POX_DIR  
$ sudo ./pox.py openflow.of_01 --port=6655  
pox.forwarding.l2_learning pox.forwarding.L3Firewall  
--l2config="l2firewall.config" --l3config="l3firewall.config"
```

Here POX controller is invoked by ./pox.py command and we run l2 learning.py application and L3Firewall.py file from POX controller. All the forwarding applications have to be stored in /pox/pox/forwarding directory. To give a relative path from the directory where the POX binary is present, we follow convention pox.forwarding.L3Firewall

4. Run mininet using containernet. Now, you can open a new terminal window, and run the following command:

```
$ sudo mn --topo=single,9 --controller=remote,port=6633  
--controller=remote,port=6655 --switch=ovsk --mac
```

It will create a mininet environment in containernet with 9 containernet hosts, one OVS switch, and two remote controllers. Option –mac will assign a small, unique, and fixed set of Mac addresses based on host ID. It will remain constant after every run.

We have started the POX controller earlier, and thus, POX will connect to this mininet environment as a remote controller. The mininet topology should look like Figure CS-CNS-00101.3.

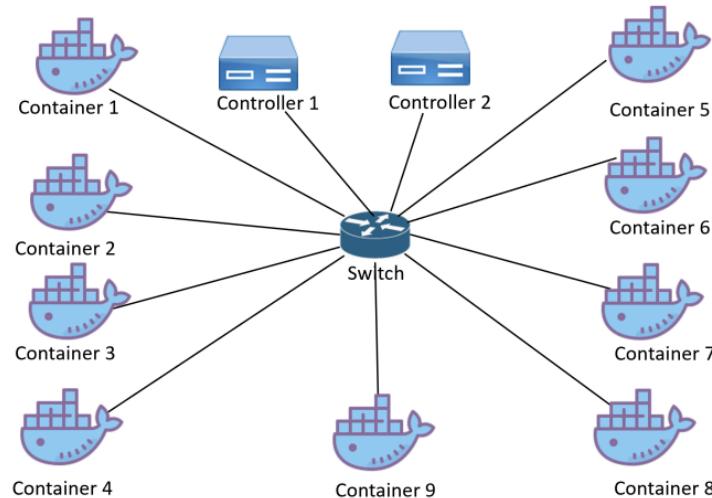


Figure CS-CNS-00101.3
Network topology for SDN.

The following table CS-CNS-00101.2 depicts the mapping between MAC addresses and respective IP addresses in specified mininet topology. Since we have enabled –mac option in mininet topology, the MAC address will be fixed for each containernet host in every run.

Container Host	Layer 2 Address (MAC address)	Layer 3 Address (IP address)
Container host h1	00:00:00:00:00:01	10.0.0.1
Container host h2	00:00:00:00:00:02	10.0.0.2
Container host h3	00:00:00:00:00:03	10.0.0.3
Container host h4	00:00:00:00:00:04	10.0.0.4
Container host h5	00:00:00:00:00:05	10.0.0.5
Container host h6	00:00:00:00:00:06	10.0.0.6
Container host h7	00:00:00:00:00:07	10.0.0.7
Container host h8	00:00:00:00:00:08	10.0.0.8
Container host h9	00:00:00:00:00:09	10.0.0.9

Table CS-CNS-00101.2

Mapping of MAC addresses with IP addresses in respective containernet hosts

Note that in the following subsections, you can use two ways to test your virtual networks from the command line of a host:

1. You can specify from which host to initiate a command. For example:

```
containernet> h1 ping h3 % ping from h1 to h3  
containernet> h1 hping3 -c 5 h2 -V --tcp-timestamp # use  
hping3 to sent tcp packets to h2
```

2. you can start a host terminal and use network configuration to perform the test. For example, the following command is to start an x-terminal of host h1:

```
containernet> xterm h1
```

Once the x-terminal is started you can run the above commands just like in a real host command line, such as:

```
$ ping 10.0.0.3 % ping from h1 to h3 (10.0.0.3)  
$ hping3 -c 5 10.0.0.2 -V --tcp-timestamp # use hping3 to sent tcp  
packets to h2 (10.0.0.2)
```

Task 2.2 Verifying the working of the firewall

In this project, Open vSwitch is used. The new rules added based on the match and controller signals can be viewed from OVS switches. To do so, you have to dump flow entries at OVS switches and use OpenvSwitch commands to verify and check the rules getting added in the Switches.

When you run L3Firewall application from POX and create a mininet environment, an SDN is formed with mininet containers as hosts, OVS as a data plane switch to handle switching functionalities, and POX to handle control plane functionality.

Open another terminal to run OVS commands that will verify the functionality of POX.

1. Check openvswitch database configuration

```
$ sudo ovs-vsctl show
```

2. list ports for specific switch

```
$ sudo ovs-vsctl list-ports s1
```

3. Check OpenFlow switch details

```
$ sudo ovs-ofctl show s1
```

4. Check OpenFlow flow details

```
$ sudo ovs-ofctl dump-flows s1
```

We have provisioned a virtual switch called OVS using mininet. ovs-ofctl provides a command line tool to manage and control OpenFlow switches. ovs-ofctl prints the output of all the flows in the same order the switch sends them.

Consider the following example output of an OpenFlow entry:

```
cookie=0x0, duration=125.376s, table=0, n_packets=4, n_bytes=392,  
priority=65535,dl_src=00:00:00:00:00:02,dl_dst=00:00:00:00:00:04 actions=drop
```

The output consists of a few important parameters which are essential in our case.

- **cookie:** This field can be added while adding or deleting a flow entry to the OpenFlow rules. We do not require this field to be updated, hence it uses the default value as 0x0.
- **Duration:** This includes the time in seconds in which the flow entry resides in the table.
- **table:** We can specify the table number here as any number between 0 to 255.
- **n packets:** This specifies the number of packets that match with this flow entry. In this particular case, a total of 4 packets matched with this entry.
- **n bytes:** This specifies the number of bytes that match with this flow entry. In this particular case, a total of 392 bytes matched with this entry.
- **priority:** Priority is a very important entry in the case of flows. A higher value will match before a lower value. Priority values can be any value between 0 to 65535. If we want a rule to get precedence in case there could be multiple matches for a packet, then the packet will follow the rule that has a higher priority value. If this value is not specified, the default value would be 32768.
- **dl src:** These are the matching criteria provided in a rule file of the firewall (either l2firewall.config or l3firewall.config). This specifies the source mac address.
- **bf dl dst:** These are the matching criteria provided in a rule file of the firewall (either l2firewall.config or l3firewall.config). This specifies the destination mac address.
- **action:** This decides what to do with a packet that matches this specified rule.

After sending the packets, if the incoming packet matches with any of the specified flows, the packet count in that specific flow entry will increase. This indicates that there is a packet matching with the criteria specified in the OpenFlow rule. This is how we know if the rules we want to install, are working in OpenFlow switch.

Note that these flows have a timeout mechanism, hence a particular flow will not be seen all the time.
In the containernet terminal, send ICMP packets from hosts to check the firewall functionality.
(in containernet terminal:)

```
containernet> h1 ping h2
```

This will send ICMP requests from containernet host h1 to containernet host h2. Since we have a rule in l3firewall.config file for blocking traffic from IP address 10.0.0.1 (this IP address belongs to containernet host h1) to destination IP address 10.0.0.2 (this IP address belongs to containernet host h2), it will add the rule from config file to OpenFlow rule in the switch.

This addition of rule can be verified from OVS commands which will dump flows at switch s1. As you start a ping from host h1 to h2, in OVS terminal you can check the rule gets added with the source address and destination address as provided in l3firewall.config file. You can also see a number of packets sent from h1 to h2.

All the packets sent and received from IP addresses are given in the l3firewall.config files are BLOCKED by the firewall application. Any other network traffic is ALLOWED by the firewall. To demonstrate this, try to ping from host h1 to host h9. There is no rule added to block network traffic from host h1 (IP: 10.0.0.1) to host h9 (10.0.0.9).

(in containernet terminal:)

```
containernet> h1 ping h9
```

Packets will be sent from container host h1 to h9. To check this, run OVS command in ovs-terminal:

```
$ sudo ovs-ofctl dump-flows s1
```

To verify the firewall rule for blocking TCP packets specified in rule number 3, send the tcp packet using the hping3 tool. (In containernet terminal window:)

```
h1 hping3 -c 5 h2 -v --tcp-timestamp
```

5. Print port statistics

```
$ sudo ovs-ofctl dump-ports-desc s1
```

6. Check status of the flow tables

```
$ sudo ovs-ofctl dump-tables s1
```

7. Check all hidden flows for particular switch

```
$ sudo ovs-appctl bridge/dump-flows s1
```

Submission Directions for Project Deliverables

You are given an unlimited number of attempts to submit your best work. The number of attempts is given to anticipate any submission errors you may have in regards to properly submitting your best work within the deadline (e.g., accidentally submitting the wrong paper). It is not meant for you to receive multiple rounds of feedback and then one (1) final submission. Only your most recent submission will be assessed.

You must submit your SDN-Based Stateless Firewall Project deliverables in submission space in your courses. Learners may not email or use other means to submit any assignment or project for review, including feedback, and grading.

The SDN-Based Stateless Firewall Project includes two (2) deliverables:

- **Project Report :** Use the Project Report template to submit your report as a DOC or PDF. Your file should be named using the following format: [Your Name_CSE 548_Name of Project or Assignment]
- **Video Demonstration:** Submit a .zip file of your video demonstration. The video should be 10 minutes or less demonstrating your implementation. Your file should be named using the following format: [Your Name_CSE 548_Name of Project or Assignment]

Project Report and Video Demonstration Submission

1. This submission will include two file uploads.
2. Navigate to your course and click on the “**Submission: SDN-Based Stateless Firewall Project**” submission space
3. Click “**Start Attempt**” in the upper right corner
4. Click “**Upload File**”
5. Locate and select your report file.
6. Click “**+ Add Another File.**” Locate and select your video demonstration file.
 - a. Upload your video demonstration as a .zip file.
7. When finished, click “**Submit Assignment**”
8. If needed: to resubmit the assignment
 - a. Click “**New Attempt**” and add your file submission again

- b. When finished, click “**Submit Assignment**”

Evaluation

You will be evaluated using the following criteria.

1. Create a mininet-based topology with 4 container hosts and one controller switch and run it.
 - Add a link from controller 1 to switch 1.
 - Add a link from controller 2 to switch 1.
 - Add a link from switch 1 to container 1.
 - Add a link from switch 1 to container 2.
 - Add a link from switch 1 to container 3.
 - Add a link from switch 1 to container 4.
2. Make the interfaces up and assign IP addresses to interfaces of container hosts.
 - Assign IP address 192.168.2.10 to container host #1.
 - Assign IP address 192.168.2.20 to container host #2.
 - Assign IP address 192.168.2.30 to container host #3.
 - Assign IP address 192.168.2.40 to container host #4.
3. Add new rule to l3config file for blocking ICMP traffic from source IP 192.168.2.10 and desti-nation IP 192.168.2.30.
4. Add new rule to l3config file for blocking ICMP traffic from source IP 192.168.2.20 and desti-nation IP 192.168.2.40.
5. Add new rule to l3config file for blocking HTTP traffic from source IP 192.168.2.20.
6. Add new rule to l2config file for blocking traffic from MAC address 00:00:00:00:00:02 to desti-nation MAC address 00:00:00:00:00:04.
7. Add new rule to l3config file for blocking tcp traffic from 192.168.2.10 to 192.168.2.20.
8. Add new rule to l3config file for blocking udp traffic from 192.168.2.10 to 192.168.2.20.
9. Important: You must submit a demo video of your work to showcase your implementation. The video should be:
 - Maximum 10 minutes long.
 - Shows successful execution of Tasks 1-8.

- Does not show your network setup, just these tasks.
- Submit the video link along with your Project report. To submit more than one file in Canvas this video file must be a .zip.